

EE324 Assignment 2

<Implementing Redis Web Server>

2020.09.29

Goal of the Assignment 2

- Implement Web Server that supports one of the widely used applications, Redis
- Improve the performance of your Web Server via multi-threads and event-based IO multiplexing

Assignment Due

- Part 1: Implement Redis Web Server **(Due ~10/13 11:59 PM)**
- Part 2: Improve Part 1 w/ multi-threads **(Due ~10/27 11:59 PM)**
- Part 3: Improve Part 1 w/ libevent **(Due ~10/27 11:59 PM)**
- **10% late penalty per day**
- **Hard deadline: 23:59, Oct. 30. 2020 (Friday)**

Assignment Due

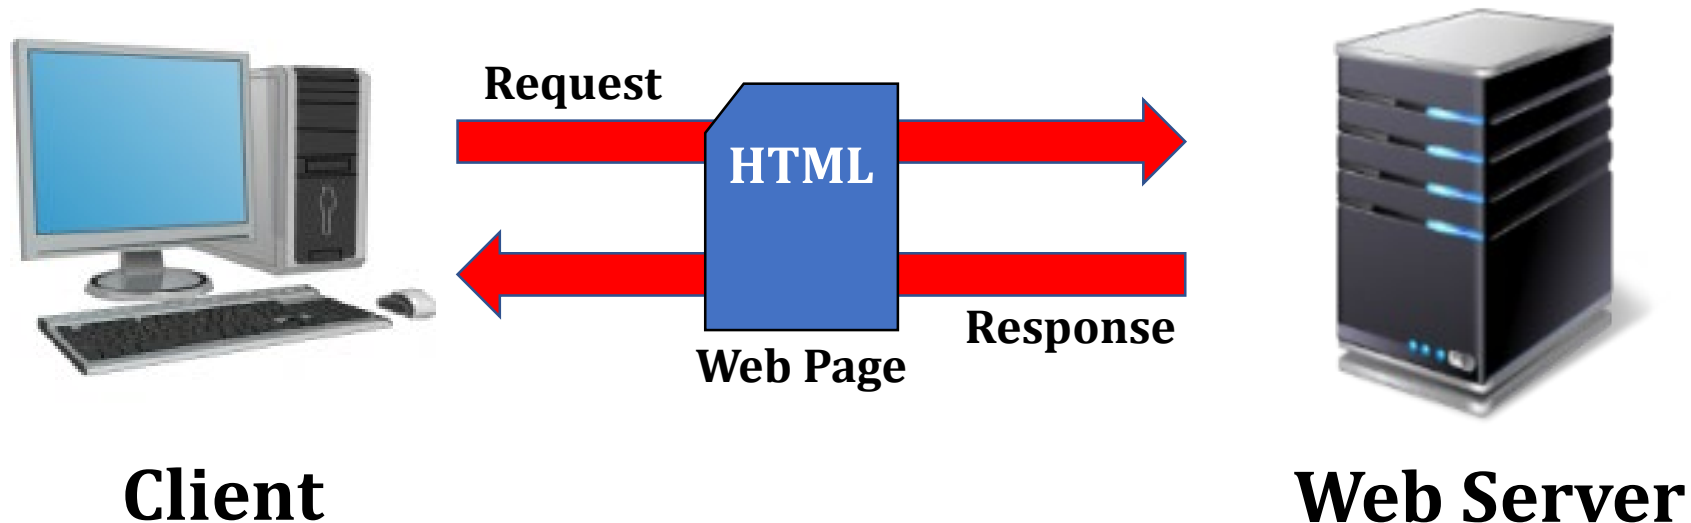
OCTOBER 2020						
Sun	Mon	Tue	Wed	Thu	Fri	Sat
27	28	29 HW2 Release	30	1	2	3
4	5	6	7	8	9	10
11	12	13 Part 1 Due	14	15	16	17
18	19	20 Mid-term Exam	21	22	23	24
25	26	27 Part 2, 3 Due	28	29	30	31

Background

- What is Web Server ?
- HTTP protocol
- What is Redis ?
- Redis Protocol

What is Web Server ?

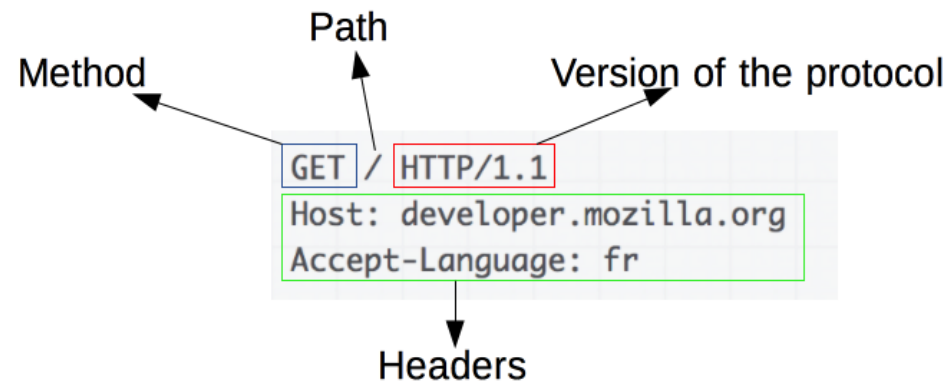
- Stores, processes and delivers web pages to clients.
- Processes incoming network requests from clients over HTTP protocol.



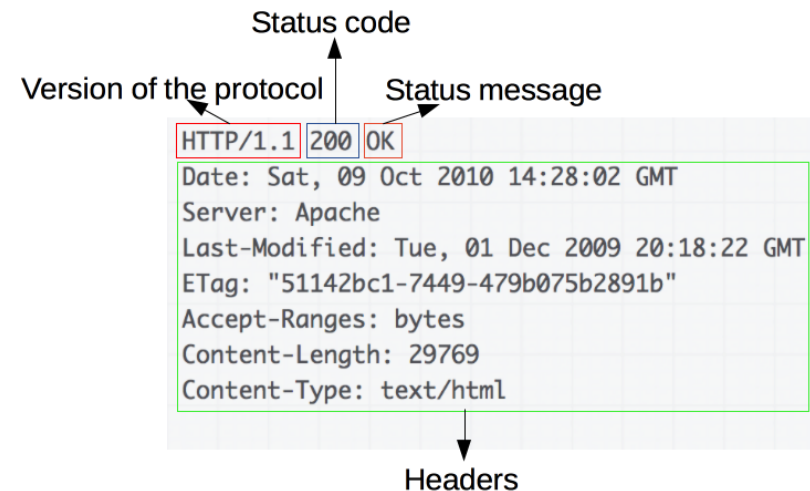
HTTP Protocol

- Defines how your client requests resources from a web server and how the server responds.
- Methods
 - GET : retrieve information identified by the request-url
 - POST : submit data to the server for updates

<HTTP GET Request Example>

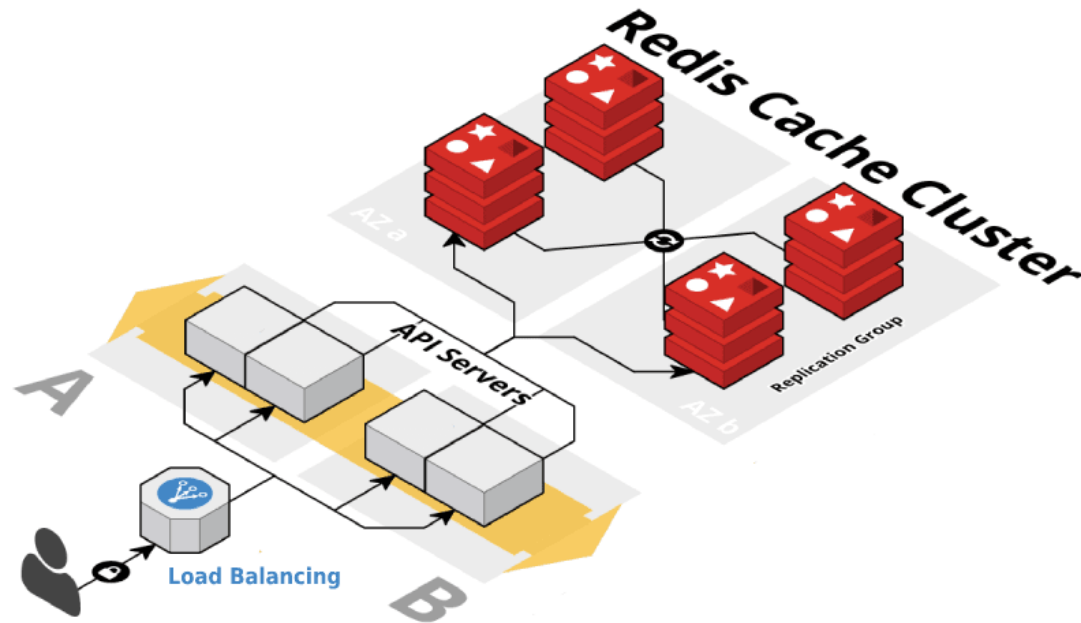


<HTTP GET Response Example>



What is Redis ? redis

- In-memory key-value data store, widely used as database and cache.
- Redis avoids seek time delays and access data in micro-seconds.
- Enables real-time apps such as Gaming, IOT, ...



Redis Example Usage

- Commands
 - SET [key] [value] : set the string value of a key
 - GET [key] : get the value of a key
 - FLUSHALL : remove all keys from database

```
jaykim305@iris1:~/redis-6.0.6/src$ ./redis-server --port 6398
6:C 28 Sep 2020 10:46:20.151 # o000o000o000o Redis is starting
6:C 28 Sep 2020 10:46:20.151 # Redis version=6.0.6, bits=64, c
6:C 28 Sep 2020 10:46:20.151 # Configuration loaded
6:M 28 Sep 2020 10:46:20.151 * Increased maximum number of ope

Redis 6.0.6 (00000000/0) 64
Running in standalone mode
Port: 6398
PID: 3717646
```

<Redis server>

```
(base) jaykim305@iris1:~/redis-6.0.6/src$ ./redis-cli -p 6398
127.0.0.1:6398> ping
PONG
127.0.0.1:6398> GET course
(nil)
127.0.0.1:6398> SET course EE324
OK
127.0.0.1:6398> SET assign2 Redis_Web_Server
OK
127.0.0.1:6398> GET course
"EE324"
127.0.0.1:6398> GET assign2
"Redis_Web_Server"
127.0.0.1:6398> FLUSHALL
OK
127.0.0.1:6398> GET course
(nil)
127.0.0.1:6398>
```

<Redis command line interface>

Redis Protocol redis

For more info, refer to : <https://redis.io/topics/protocol>

<Rules>

- *: followed by the number of elements in the command array
- \$: followed by the number of bytes composing the string, terminated by CRLF(\r\n)

<Example>

GET course

*2\r\n

\$3\r\n

GET\r\n

\$6\r\n

course\r\n

SET course EE324

*3\r\n

\$3\r\n

SET\r\n

\$6\r\n

course\r\n

\$5\r\n

EE324\r\n

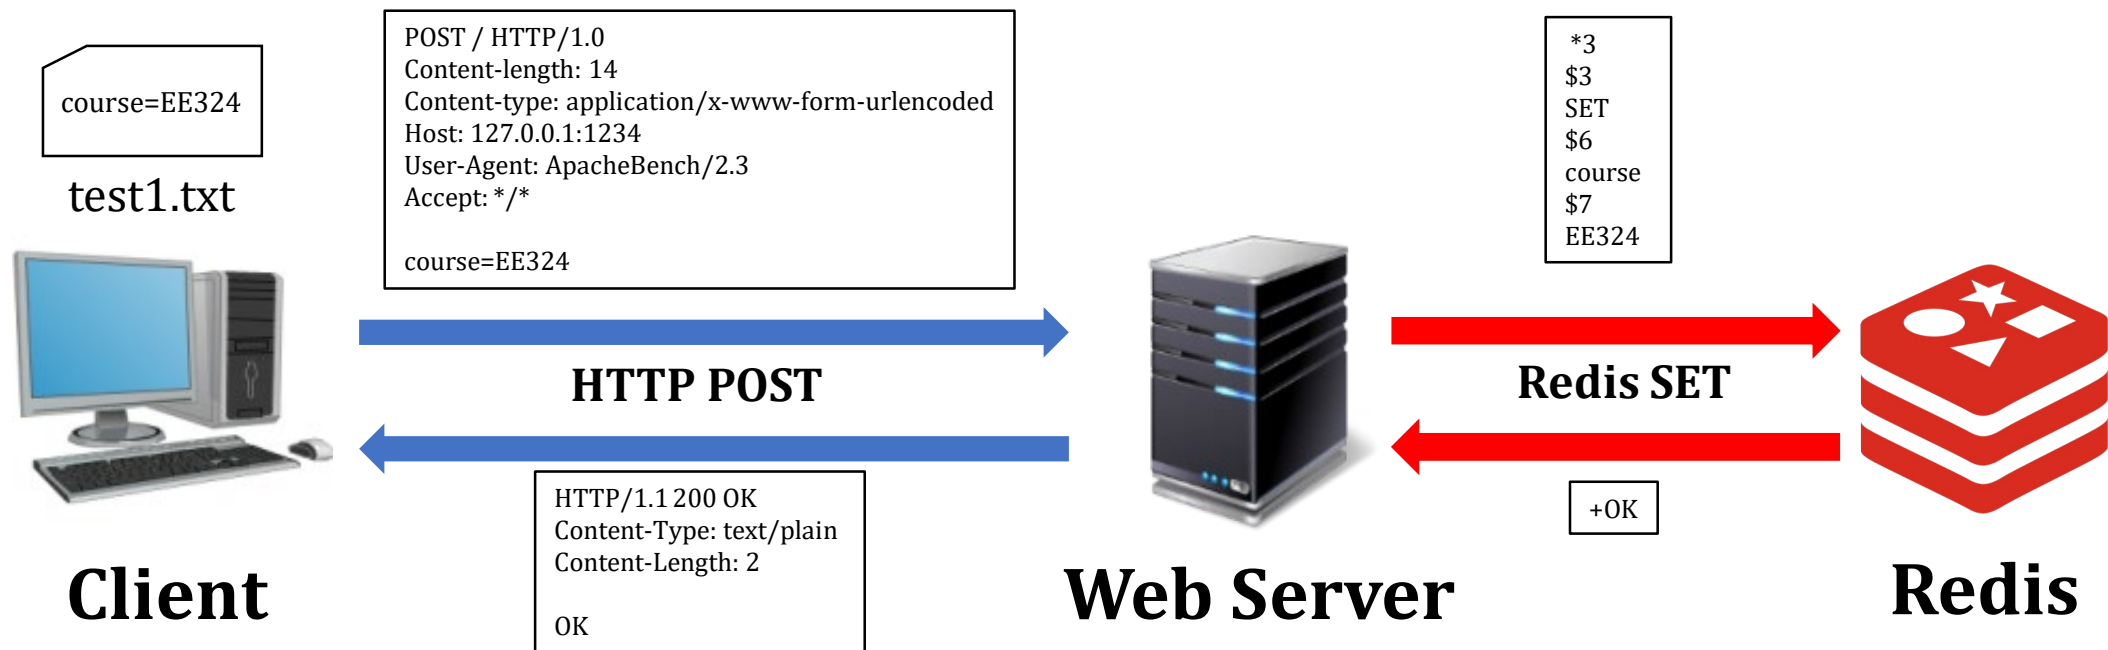
Part 1: Implement Redis Web Server

- You need to implement Web Server that communicates with Redis.
 - Client <---- HTTP Protocol (GET, POST) ----> Web Server
 - Web Server <---- Redis Protocol (GET, SET) ----> Redis
 - Should handle multiple concurrent clients : use fork () as used in HW1



Scenario1 : SET key-value via HTTP POST

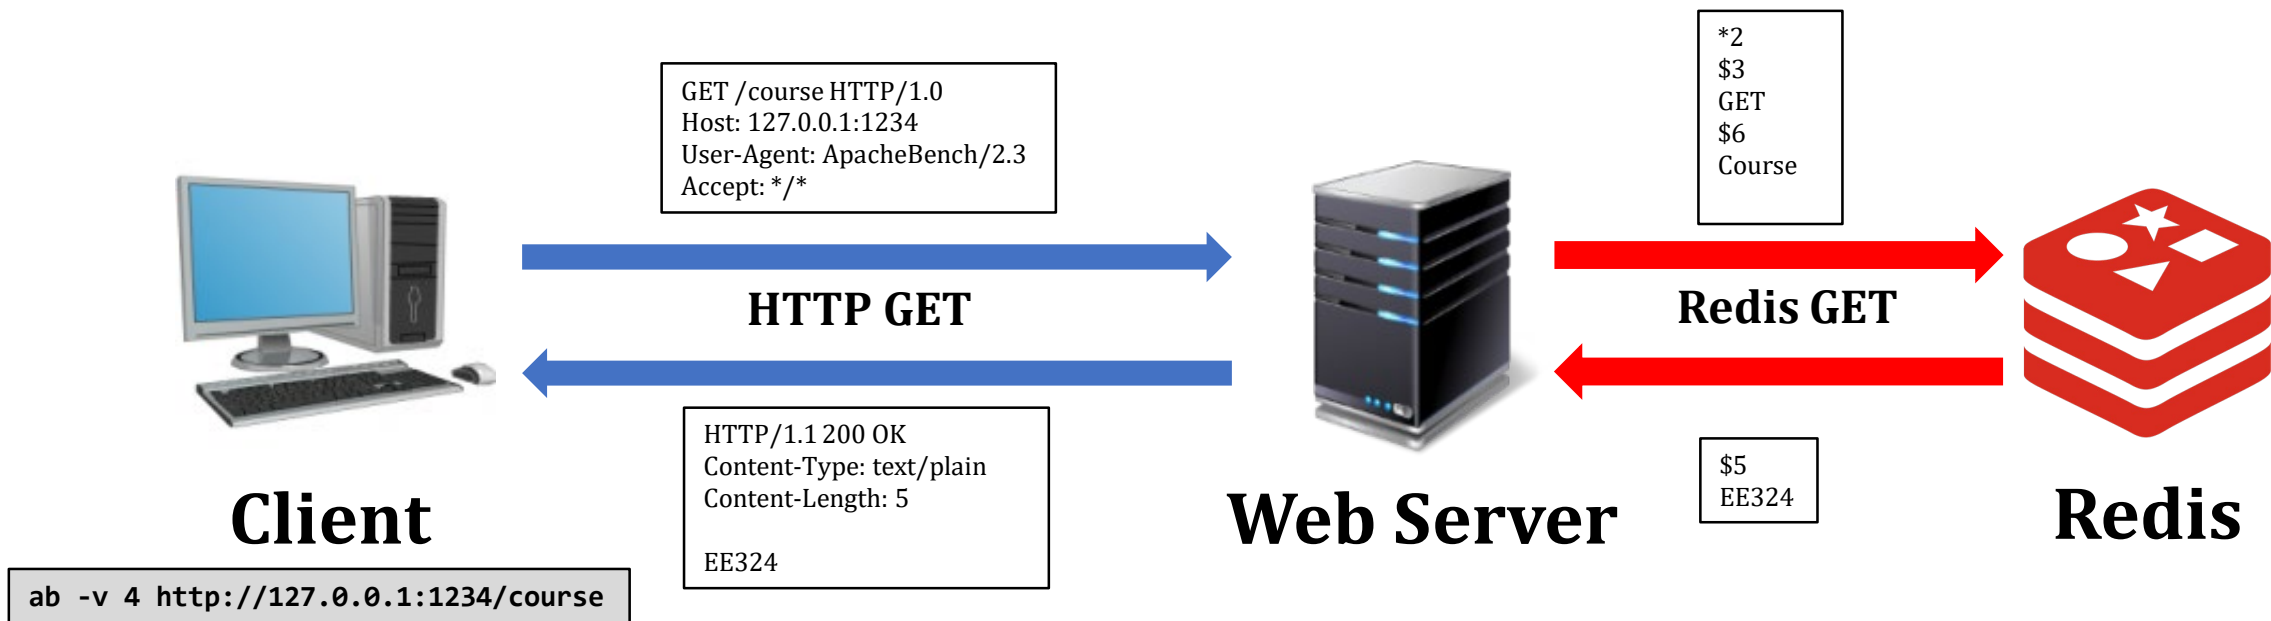
- Client wants to store “course” key as a value “EE324” to the server



```
ab -v 4 -p test1.txt http://127.0.0.1:1234/
```

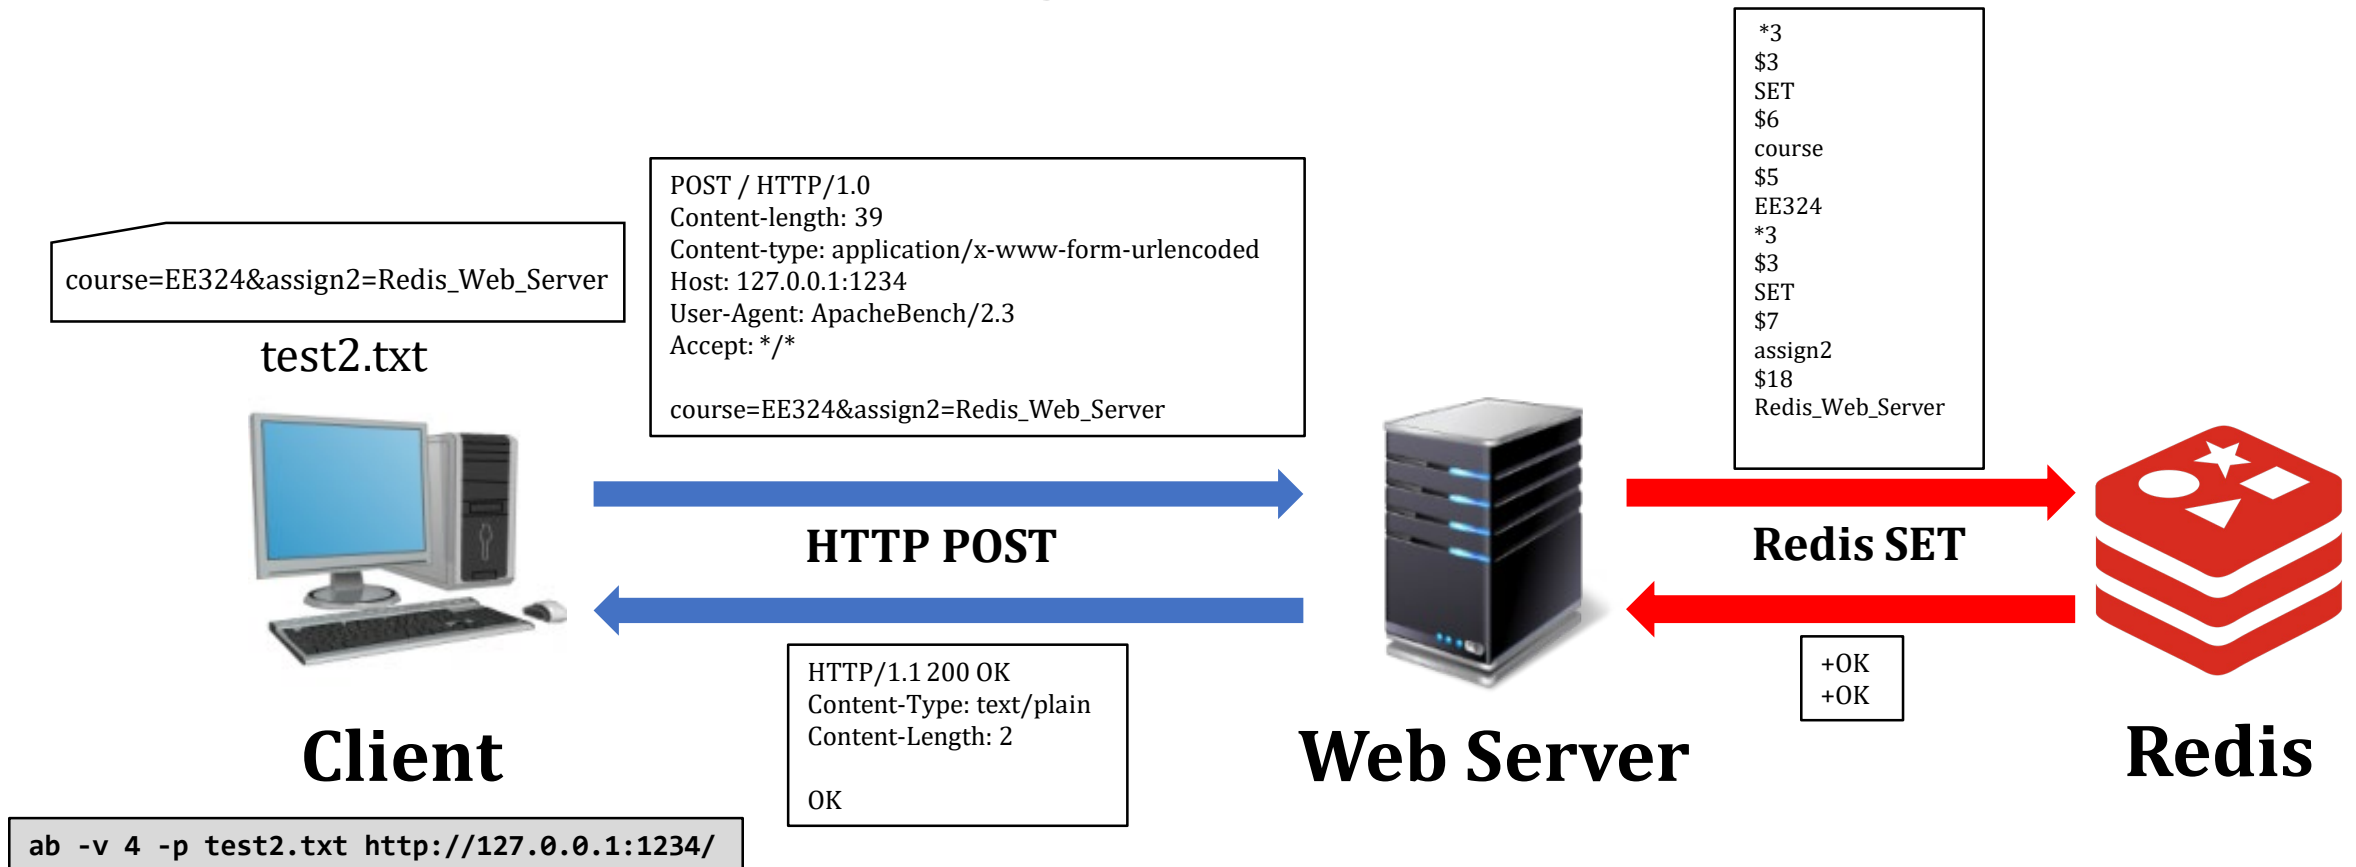
Scenario2 : GET key-value via HTTP GET

- Client wants to get the value of “course” from the server



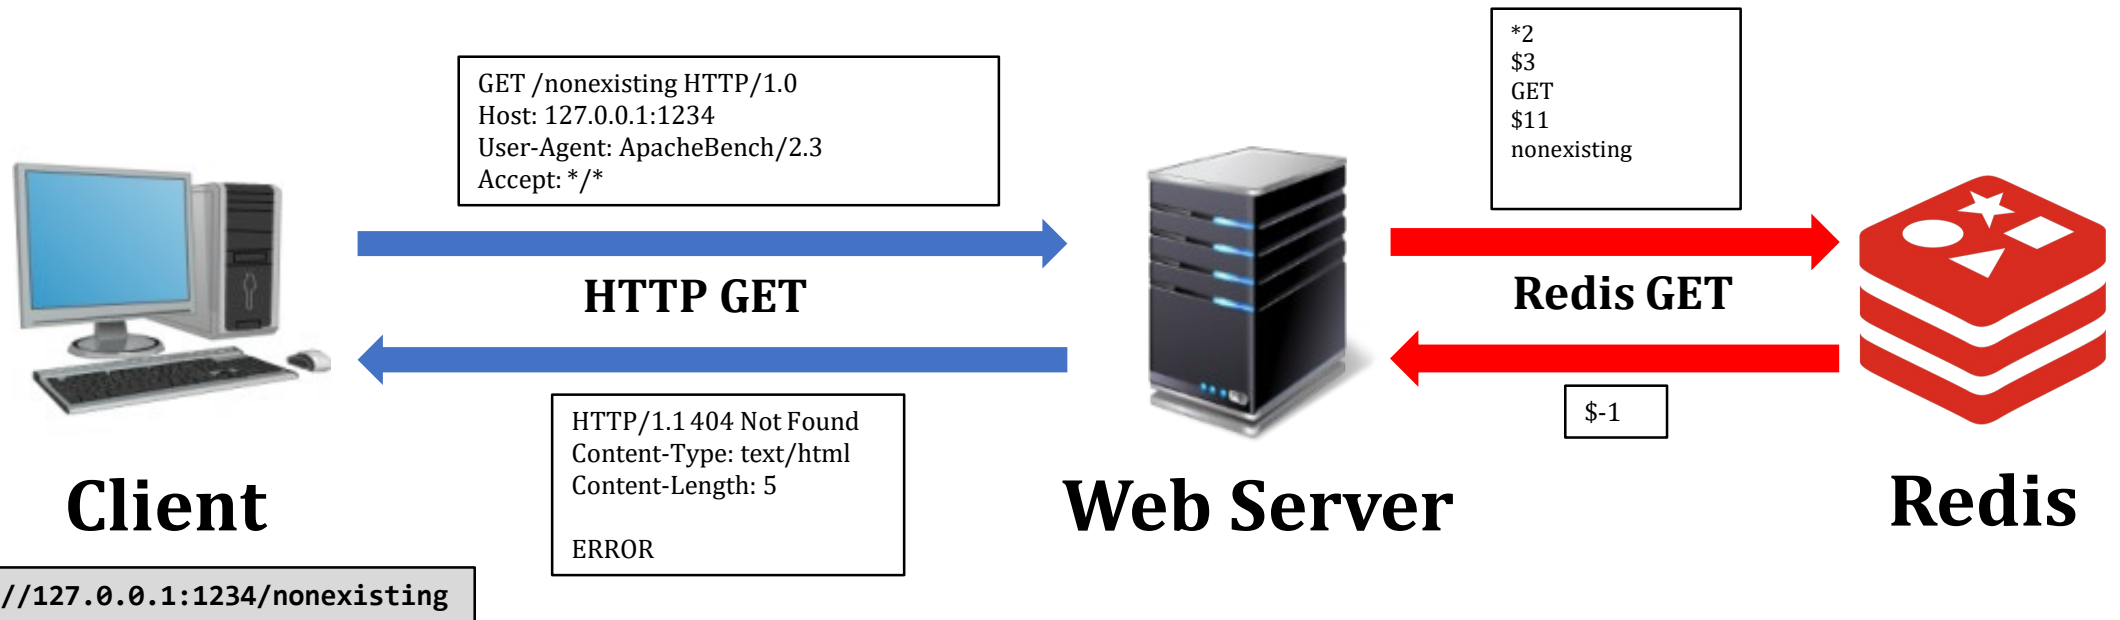
Scenario3: Multiple SET key-value

- Client wants to store multiple key-value pairs to the server
{course=EE324} and {assign2=Redis_Web_Server}



Scenario4: GET non-existing key

- Client tries to get value of a key that is not stored



Required Functionality (1/2)

- Your Web Server should
 1. parse HTTP request (GET, POST).
 2. get the corresponding Redis command (i.e., GET, SET) and its key-value arguments.
 3. build Redis request following Redis protocol.
 4. retrieve Redis response from Redis Server.
 5. send response to client via HTTP response body.
 - handle error cases (non-existing keys, invalid Redis protocol)
 - single(multiple) key-value SET and single key-value GET
- Your Web Server should support multiple concurrent clients (up to 1K).
 - part 1 : fork, part 2: multi-threads, part 3: libevent
- Your server must at least handle key and value of any size.

Required Functionality (2/2)

- Your server binaries (webserver_fork, webserver_thread, webserver_libevent) should get
[Port # of your Web Server], [IP address of your Redis Server],
[Port # of your Redis Server] from stdin.

- Example)

```
bin/webserver_fork 1234 127.0.0.1 6398
```

Assignment Due

- Part 1: Implement Redis Web Server (fork based)
(Due ~10/13 11:59 PM)
- Part 2: Improve Part 1 w/ multi-threads (Due ~10/27 11:59 PM)
- Part 3: Improve Part 1 w/ libevent (Due ~10/27 11:59 PM)
- 10% late penalty per day
- Hard deadline: 23:59, Oct. 30. 2020 (Friday)

Part 2: Improve Part 1 w/ multi-threads

- Threads are light-weight process and have less overhead than fork.
- Use `pthread.h`
- thread pool or using select is not required (this is for HW3)
- Simply spawn threads in a similar way as `fork()`
- Use
 `pthread_create()`, `pthread_join()`, `pthread_detach()`
- Compare the performance of Part 2 and 1 using Apache Bench (ab).

Part 3: Improve Part 1 w/ libevent

- Callback based event interface based on event mechanisms such as epoll, kqueue, select and poll.
- Use `event2/event.h`
- The whole process runs on single main thread
- Use
`event_base_new()`, `event_base_dispatch()`,
`event_new()`, `event_del()`, `event_add()`
- Compare the performance of Part 3 and 1 using Apache Bench (ab).

Assignment Due

- Part 1: Implement Redis Web Server (Due ~10/13 11:59 PM)
- Part 2: Improve Part 1 w/ multi-threads (Due ~10/27 11:59 PM)
- Part 3: Improve Part 1 w/ libevent (Due ~10/27 11:59 PM)
- **10% late penalty per day**
- **Hard deadline: 23:59, Oct. 30. 2020 (Friday)**

Before you start

- Install Redis (<https://redis.io/download>)

```
$ wget http://download.redis.io/releases/redis-6.0.8.tar.gz  
$ tar xzf redis-6.0.8.tar.gz  
$ cd redis-6.0.8  
$ make
```


- Your `redis-server`, `redis-cli` binaries are located in `src`.
- Install `libevent` (this is not required for eelab machines)
`sudo apt-get install libevent-dev`
- Download template from gitlab (Clone repository)
`https://gitlab.ee324.kaist.ac.kr/assignments/assignment-2.git`



Grading

1. Part 1 : (50 pts)
 - Functionality – 25 pts
 - Robustness – 25 pts
2. Part 2 : (10 pts)
 - Functionality – 10 pts
3. Part 3: (40 pts)
 - Functionality – 10 pts
 - Robustness – 10 pts
 - Performance – 20 pts
4. Bonus points (5 pts)

(+5pt) Bonus: Percent-encoding

- Key and value sent from client is URL encoded (percent-encoding)
 - [Content-type: application/x-www-form-urlencoded](https://en.wikipedia.org/wiki/Percent-encoding)
 - <https://en.wikipedia.org/wiki/Percent-encoding>

hello, world!  hello%2C%20world%21

  %F0%9F%8D%BA

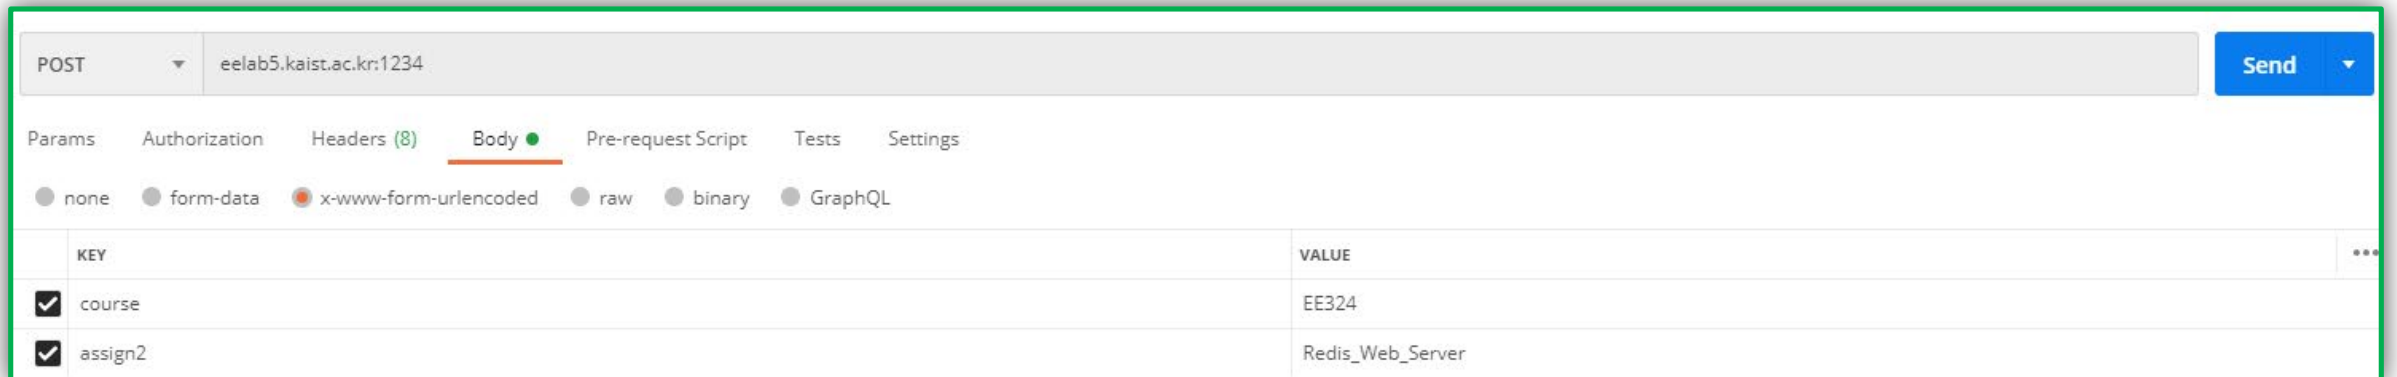
- Decode the request when you GET/SET value on Redis
- Example (SET)
 - Input from HTTP client: R%25D=staff%40ee324
 - Parsed key and value: {R&D, staff@ee324}
 - You should not store “R%25D” as is!
- **Note decoded data can be binary!**
- You can use helper function in `src/urllencode.c` (if you want.)

Test Cases

- We provide `test.py`
- We will use more tricky test cases not included in the `test.py`
- Usage:
 - `test/test.py --redis-port [port] --test [test case] [binary]`
 - `test case = {func_get, func_set_simple, func_set_multiple}`
 - `binary = {bin/webserver_fork, bin/webserver_thread, bin/webserver_libevent}`

Useful Tools for Developing

- ApacheBench (ab)
 - We will use ApacheBench (ab) as a benchmark to test the performance of your web servers.
- Postman (<https://www.postman.com/downloads/>)



Performance Test using Apache Bench (ab)

1. ab with 1 concurrency 1,000 requests
 - `$ ab -n 1000 server_ip:server_port`
 2. Ab with 100 concurrency, 10,000 requests
 - `$ ab -c 100 -n 10000 server_ip:server_port`
- We will test your program up to 1K requests with 1K concurrency.
 - Other options (type ab for more options)
 - p : postfile
 - v : verbosity
 - T : content-type
- (Note for Bonus: you should set this to [application/x-www-form-urlencoded](#))

Running Example

POST eelab5.kaist.ac.kr:1234 **Postman Application** Send

Params Authorization Headers (8) **Body** Pre-request Script Tests Settings

none form-data x-www-form-urlencoded raw binary GraphQL

KEY	VALUE
<input checked="" type="checkbox"/> course	EE324
<input checked="" type="checkbox"/> assign2	Redis_Web_Server

```
ab -v 4 -p test2.txt http://eelab5.kaist.ac.kr:1234/
```

Apache Bench

```
./webserver_fork 1234 127.0.0.1 6398
```



Client



HTTP Protocol



**Web Server
(eelab5)**



Redis Protocol



Redis

```
./redis-server --port 6398
```

References

- libevent Document:
<http://www.wangafu.net/~nickm/libevent-book/>,
http://www.wangafu.net/~nickm/libevent-2.0/doxygen/html/dir_db160b4728e6067cf5f9cc14ec42c79d.html
- EE324 Assignment 2 Document :
<https://www.notion.so/jaykim305/EE324-Assignment-2-Implementing-Redis-Web-Server-db47f65642454e749b26d7c0e174272f>

Submission (1/2)

- You should upload your code on GitLab
- We will download your project when the due is over
- In the project, there should be following files
 - README.md
 - Makefile
 - src/webserver_fork.c
 - src/webserver_thread.c
 - src/webserver_libevent.c
 - and some test scripts...

Submission (2/2)

- You must include the following contents in the README.md
 - the date and time of your submission in README.md
 - Performance measurements from Part2 (not Due on Part1)
 - Performance measurements from Part3 (not Due on Part1)
 - Brief explanation and reasonings about your results
- Feel free to change Makefile, test.py, or gitlab.ci.yml
(As long as your Makefile generates binaries with designated names)

Performance measurements example

```
Document Path: /ee324
Document Length: 8002 bytes ← payload size (8kB)

Concurrency Level: 1000
Time taken for tests: 8.780 seconds ← Time taken for tests
Complete requests: 1000 ← Completed requests
Failed requests: 0
Total transferred: 8068000 bytes
HTML transferred: 8002000 bytes
Requests per second: 113.89 [#/sec] (mean)
Time per request: 8780.228 [ms] (mean)
Time per request: 8.780 [ms] (mean, across all concurrent requests)
Transfer rate: 897.35 [Kbytes/sec] received

Connection Times (ms)
      min  mean[+/-sd] median   max
Connect:    0   297 695.7     20   3004
Processing:  2   630 879.5    220   6416
Waiting:    2   629 879.5    220   6416
Total:      2   926 1198.1   622   7417

Percentage of the requests served within a certain time (ms)
 50%    622
 66%   1100
 75%   1315
 80%   1349
 90%   2601
 95%   3414
 98%   4205
 99%   7411 ← 99%-ile completion time
100%   7417 (longest request)
```

- Fill README.md

Performance measurements

Set 8kB-sized value into a specific key. Measure the time for running 1,000 concurr

- Part 1
 - Completed requests: ____
 - Time taken for test: ____ ms
 - 99%-ile completion time: ____ ms
- Part 2
 - Completed requests: ____
 - Time taken for test: ____ ms
 - 99%-ile completion time: ____ ms
- Part 3
 - Completed requests: ____
 - Time taken for test: ____ ms
 - 99%-ile completion time: ____ ms

Briefly compare the performance of part 1 through 3 and explain the results.

Test Environment

- Language: C or C++
- Test O/S: Ubuntu 16.04 LTS (Xenial Xerus), Ubuntu 18.04.5 LTS (Bionic Beaver), 20.04 LTS (focal fossa) 64bit
- **We will not consider your compilation and execution problems due to the different OS versions.**

Plagiarism

- You can discuss with your colleagues, but you should turn in your own programs
- Copy and Paste
 - Will run plagiarism detection on source code
 - “Copy and paste” codes will get severely penalized
 - If detected, 0 point for all assignments (both providers and consumers)
 - But you will have a chance to defend yourself

Questions

- Please use Piazza for asking any questions.
- Please read the documents and the slide carefully before you ask any questions.